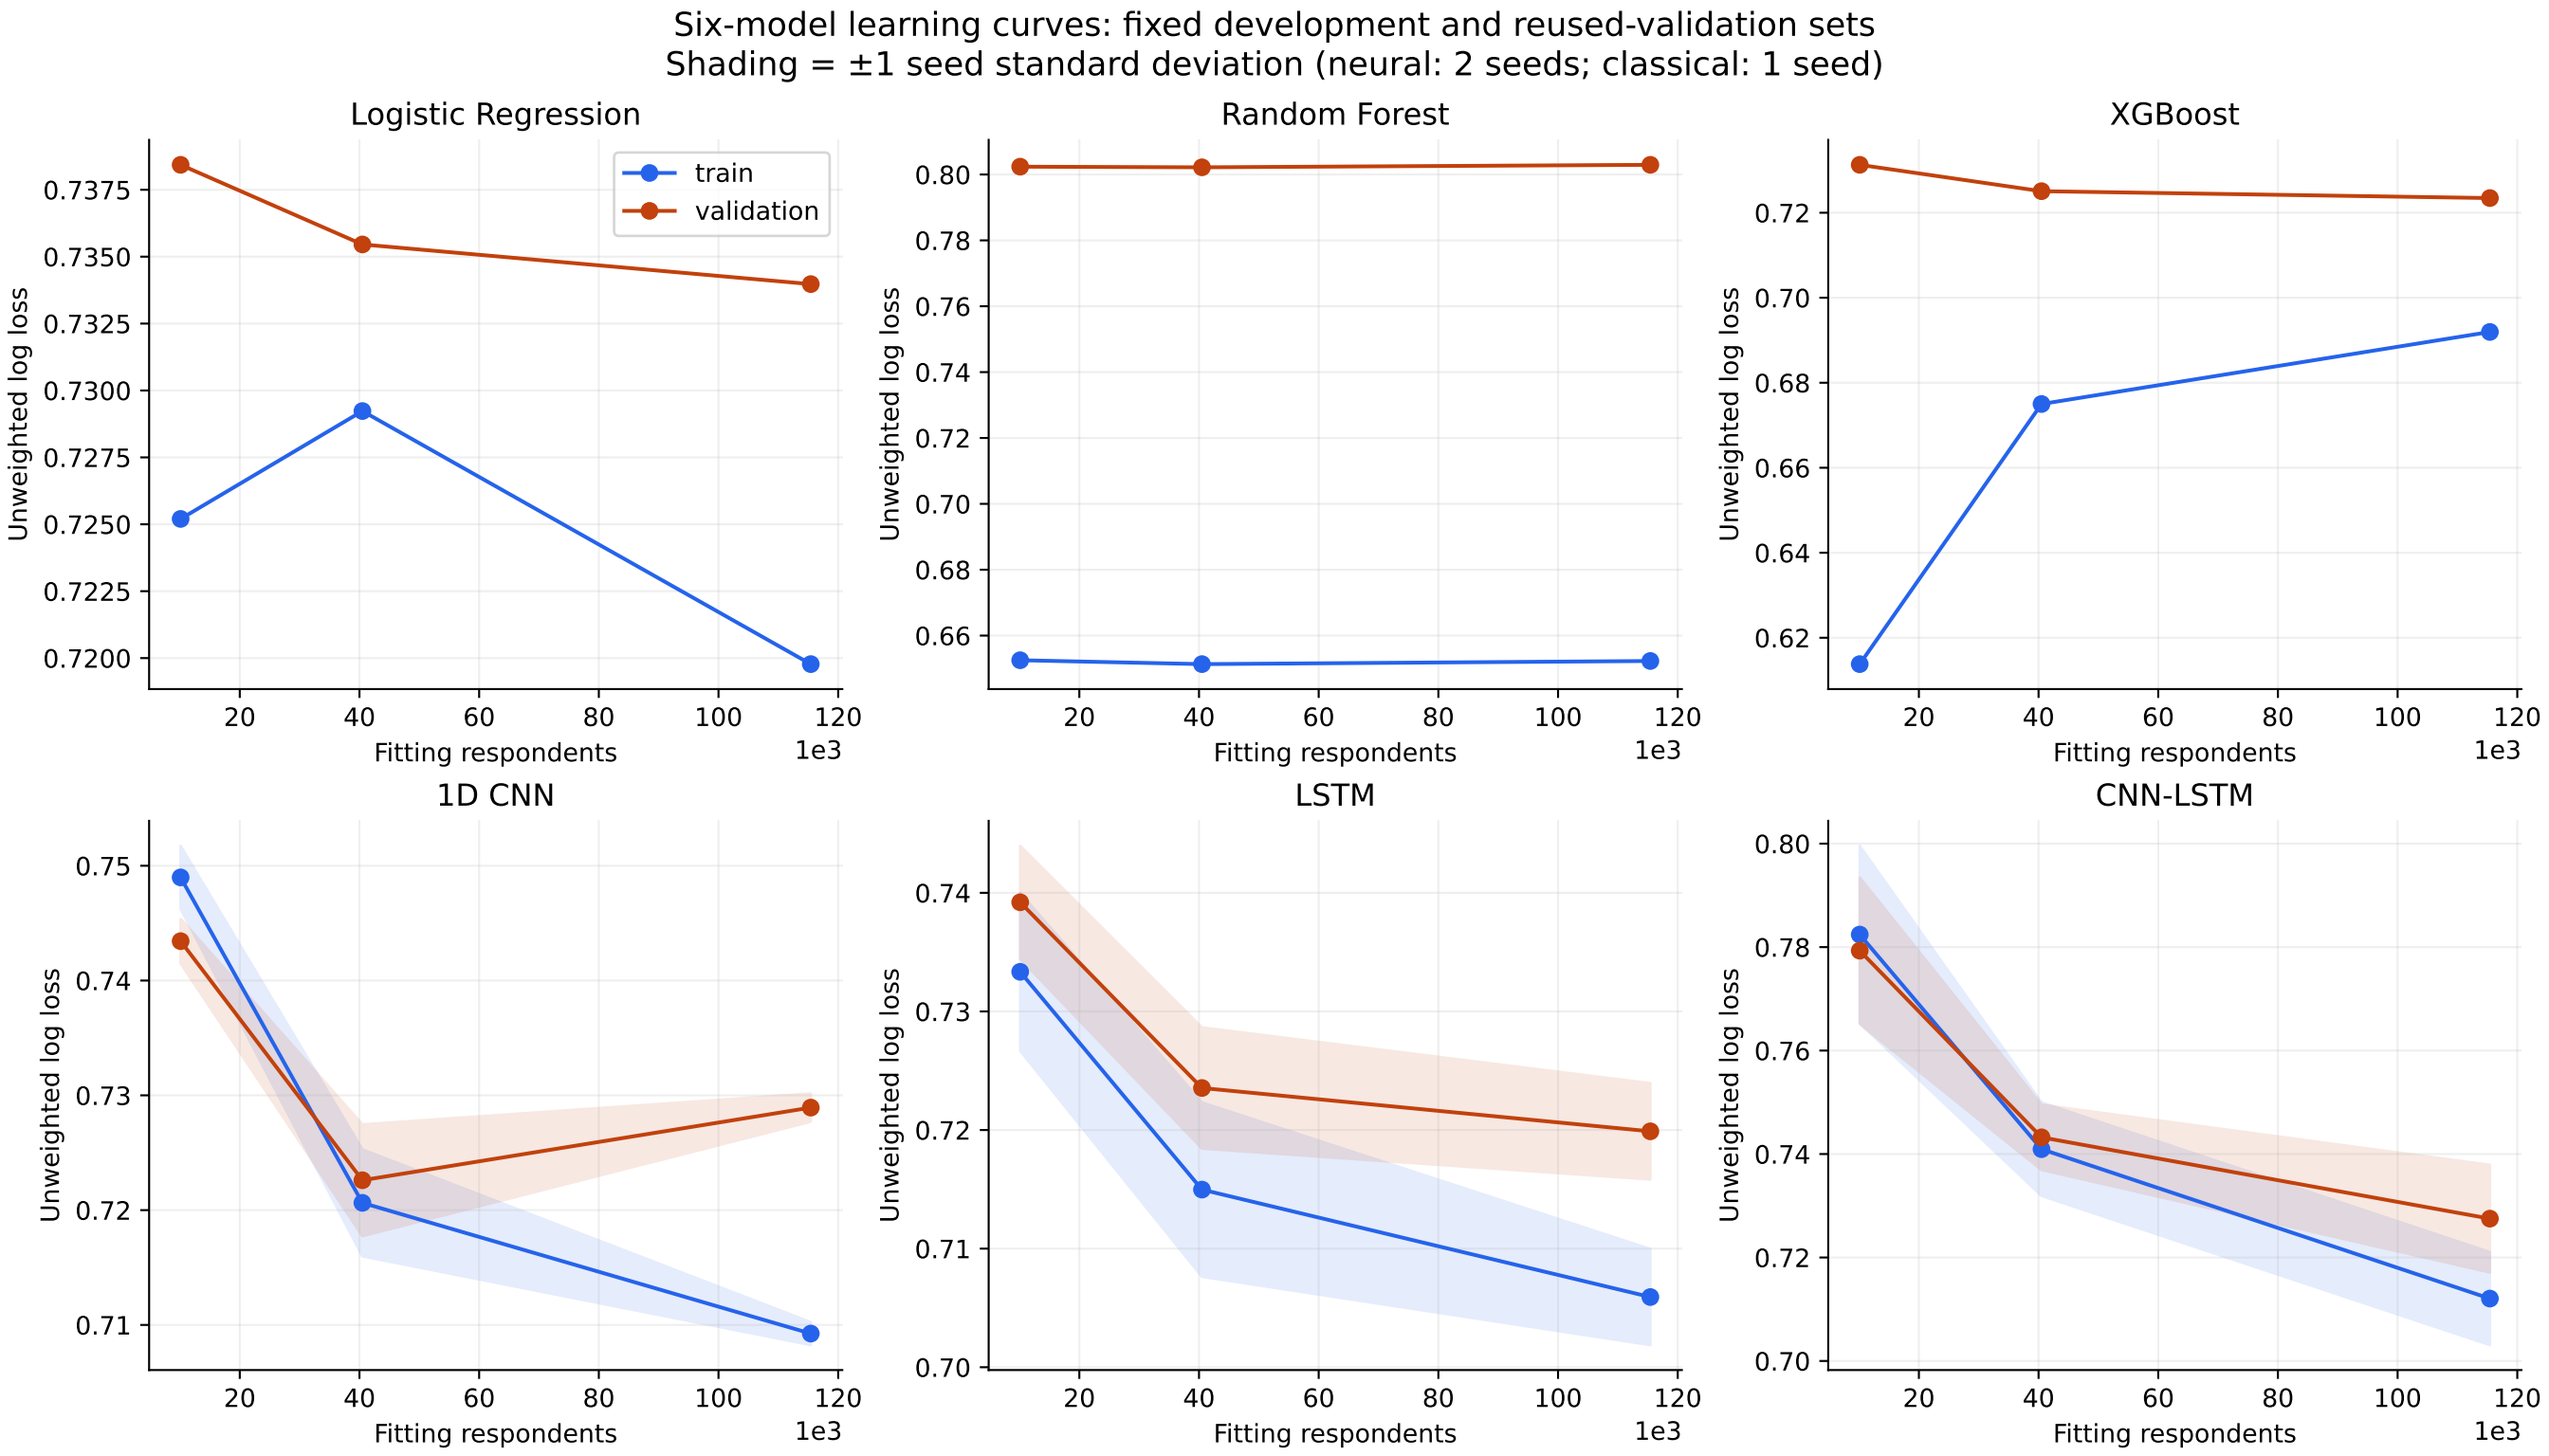
Six-model learning curves: fixed development and reused-validation sets
Shading = ±1 seed standard deviation (neural: 2 seeds; classical: 1 seed)
Logistic Regression
Random Forest
XGBoost
1D CNN
LSTM
CNN-LSTM
train
validation
Unweighted log loss
Fitting respondents
1e3